

pandas

March 21, 2022

1 PANDAS TUTORIAL (QUICK TOUR-10 MINUTES)

This is a short introduction to **pandas**, geared mainly for new users.

Customarily, we import as follow:

```
[ ]: import numpy as np
import pandas as pd
```

1.1 Object Creation

Creating a **Series** by passing a list of values, letting **pandas** create a default integer index.

```
[ ]: s = pd.Series([1, 2, -7, 10, np.nan, 8, 11])
s
```

```
[ ]: 0    1.0
1    2.0
2   -7.0
3   10.0
4    NaN
5    8.0
6   11.0
dtype: float64
```

Creating a **DataFrame** by passing a NumPy array, with datetime index and labeled columns.

```
[ ]: dates = pd.date_range("20220319", periods=7)
dates
```

```
[ ]: DatetimeIndex(['2022-03-19', '2022-03-20', '2022-03-21', '2022-03-22',
                    '2022-03-23', '2022-03-24', '2022-03-25'],
                    dtype='datetime64[ns]', freq='D')
```

```
[ ]: df = pd.DataFrame(np.random.randn(7, 4), index=dates, columns=list("ABCD"))
df
```

```
[ ]:
```

	A	B	C	D
2022-03-19	0.228181	0.411725	-0.888684	-0.907902
2022-03-20	-1.676719	1.090258	-1.966077	0.909066
2022-03-21	0.244249	-0.256577	-0.099083	-0.169163
2022-03-22	1.776478	0.235238	-0.093981	0.867664
2022-03-23	-0.185076	-0.964747	-0.542850	-0.283645
2022-03-24	-0.216682	0.182571	0.885816	-0.151794
2022-03-25	0.624784	-0.974044	1.100455	0.893449

Create a **DataFrame** by passing a *dictionary* of objects that can be converted into a series-like structure.

```
[ ]: df2 = pd.DataFrame(
    {
        "A": 1.5,
        "B": pd.Timestamp("20220319"),
        "C": pd.Series(1, index=list(range(4)), dtype=np.float32),
        "D": np.array([5] * 4, dtype=np.int32),
        "E": pd.Categorical(["test", "train", "test", "train"]),
        "F": "foo",
    }
)
df2
```

```
[ ]:
```

	A	B	C	D	E	F
0	1.5	2022-03-19	1.0	5	test	foo
1	1.5	2022-03-19	1.0	5	train	foo
2	1.5	2022-03-19	1.0	5	test	foo
3	1.5	2022-03-19	1.0	5	train	foo

The columns of the resulting **DataFrame** have different **dtypes**.

```
[ ]: df2.dtypes
```

```
[ ]: A          float64
     B    datetime64[ns]
     C          float32
     D           int32
     E         category
     F           object
dtype: object
```

1.2 Viewing data

Here is how to view the top and bottom rows of the frame.

```
[ ]: df.head()
```

```
[ ]:
```

	A	B	C	D
2022-03-19	0.228181	0.411725	-0.888684	-0.907902
2022-03-20	-1.676719	1.090258	-1.966077	0.909066
2022-03-21	0.244249	-0.256577	-0.099083	-0.169163
2022-03-22	1.776478	0.235238	-0.093981	0.867664
2022-03-23	-0.185076	-0.964747	-0.542850	-0.283645

```
[ ]: df.tail()
```

```
[ ]:
```

	A	B	C	D
2022-03-21	0.244249	-0.256577	-0.099083	-0.169163
2022-03-22	1.776478	0.235238	-0.093981	0.867664
2022-03-23	-0.185076	-0.964747	-0.542850	-0.283645
2022-03-24	-0.216682	0.182571	0.885816	-0.151794
2022-03-25	0.624784	-0.974044	1.100455	0.893449

You can pass arguments to the `head()` or `tail()` functions to display a specified number of rows.

```
[ ]: df.head(2)
```

```
[ ]:
```

	A	B	C	D
2022-03-19	0.228181	0.411725	-0.888684	-0.907902
2022-03-20	-1.676719	1.090258	-1.966077	0.909066

Display the index, columns.

```
[ ]: df.index
```

```
[ ]: DatetimeIndex(['2022-03-19', '2022-03-20', '2022-03-21', '2022-03-22',
                    '2022-03-23', '2022-03-24', '2022-03-25'],
                    dtype='datetime64[ns]', freq='D')
```

```
[ ]: df.columns
```

```
[ ]: Index(['A', 'B', 'C', 'D'], dtype='object')
```

`DataFrame.to_numpy()` gives a NumPy representation of the underlying data. Note that this can be an expensive operation when your `DataFrame` has columns with *different* data types, which comes down to a fundamental difference between pandas and NumPy: NumPy arrays have one dtype for the entire array, while pandas DataFrames have one dtype per column. When you call `DataFrame.to_numpy()`, pandas will find the NumPy dtype that can hold all of the dtypes in the `DataFrame`. This may end up being object, which requires casting every value to a Python object.

For `df`, our `DataFrame` of all floating-point values, `DataFrame.to_numpy()` is fast and doesn't require copying data

```
[ ]: df.to_numpy()
```

```
[ ]: array([[ 0.22818121,  0.41172466, -0.88868392, -0.90790186],
          [-1.67671899,  1.09025807, -1.96607677,  0.90906625],
          [ 0.24424891, -0.2565773 , -0.09908294, -0.16916329],
          [ 1.77647805,  0.23523834, -0.09398143,  0.86766426],
          [-0.18507631, -0.96474687, -0.54284961, -0.28364508],
          [-0.21668238,  0.18257144,  0.88581603, -0.15179394],
          [ 0.62478439, -0.97404414,  1.10045518,  0.893449  ]])
```

For df2, the **DataFrame** with multiply dtypes, `DataFrame.to_numpy()` is relative expensive.

```
[ ]: df2.to_numpy()
```

```
[ ]: array([[1.5, Timestamp('2022-03-19 00:00:00'), 1.0, 5, 'test', 'foo'],
          [1.5, Timestamp('2022-03-19 00:00:00'), 1.0, 5, 'train', 'foo'],
          [1.5, Timestamp('2022-03-19 00:00:00'), 1.0, 5, 'test', 'foo'],
          [1.5, Timestamp('2022-03-19 00:00:00'), 1.0, 5, 'train', 'foo']],
          dtype=object)
```

`describe()` shows a quick statistic summary of your data.

```
[ ]: df.describe()
```

```
[ ]:
      count      A      B      C      D
count  7.000000  7.000000  7.000000  7.000000
mean    0.113602 -0.039368 -0.229200  0.165382
std     1.037642  0.750703  1.048530  0.723698
min    -1.676719 -0.974044 -1.966077 -0.907902
25%    -0.200879 -0.610662 -0.715767 -0.226404
50%     0.228181  0.182571 -0.099083 -0.151794
75%     0.434517  0.323481  0.395917  0.880557
max     1.776478  1.090258  1.100455  0.909066
```

```
[ ]: df2.describe()
```

```
[ ]:
      count      A      C      D
count  4.0  4.0  4.0
mean    1.5  1.0  5.0
std     0.0  0.0  0.0
min     1.5  1.0  5.0
25%     1.5  1.0  5.0
50%     1.5  1.0  5.0
75%     1.5  1.0  5.0
max     1.5  1.0  5.0
```

Transposing your data.

```
[ ]: df.T
```

```
[ ]:      2022-03-19  2022-03-20  2022-03-21  2022-03-22  2022-03-23  2022-03-24  \
A      0.228181   -1.676719    0.244249    1.776478   -0.185076   -0.216682
B      0.411725    1.090258   -0.256577    0.235238   -0.964747    0.182571
C     -0.888684   -1.966077   -0.099083   -0.093981   -0.542850    0.885816
D     -0.907902    0.909066   -0.169163    0.867664   -0.283645   -0.151794

      2022-03-25
A      0.624784
B     -0.974044
C      1.100455
D      0.893449
```

Sorting by an axis

```
[ ]: df.sort_index(axis=0, ascending=False)
```

```
[ ]:      A      B      C      D
2022-03-25  0.624784 -0.974044  1.100455  0.893449
2022-03-24 -0.216682  0.182571  0.885816 -0.151794
2022-03-23 -0.185076 -0.964747 -0.542850 -0.283645
2022-03-22  1.776478  0.235238 -0.093981  0.867664
2022-03-21  0.244249 -0.256577 -0.099083 -0.169163
2022-03-20 -1.676719  1.090258 -1.966077  0.909066
2022-03-19  0.228181  0.411725 -0.888684 -0.907902
```

```
[ ]: df.sort_index(axis=1, ascending=False)
```

```
[ ]:      D      C      B      A
2022-03-19 -0.907902 -0.888684  0.411725  0.228181
2022-03-20  0.909066 -1.966077  1.090258 -1.676719
2022-03-21 -0.169163 -0.099083 -0.256577  0.244249
2022-03-22  0.867664 -0.093981  0.235238  1.776478
2022-03-23 -0.283645 -0.542850 -0.964747 -0.185076
2022-03-24 -0.151794  0.885816  0.182571 -0.216682
2022-03-25  0.893449  1.100455 -0.974044  0.624784
```

Sorting by values.

```
[ ]: df.sort_values(by="B")
```

```
[ ]:      A      B      C      D
2022-03-25  0.624784 -0.974044  1.100455  0.893449
2022-03-23 -0.185076 -0.964747 -0.542850 -0.283645
2022-03-21  0.244249 -0.256577 -0.099083 -0.169163
2022-03-24 -0.216682  0.182571  0.885816 -0.151794
2022-03-22  1.776478  0.235238 -0.093981  0.867664
2022-03-19  0.228181  0.411725 -0.888684 -0.907902
2022-03-20 -1.676719  1.090258 -1.966077  0.909066
```

1.3 Selection

1.3.1 Getting

Selecting a single column, which yeild a **Series**, equivalent to `df.A`.

```
[ ]: print("df.A", df.A, sep="\n")
      print("="*30)
      print("df[\"A\"]", df["A"], sep="\n")
```

```
df.A
2022-03-19    0.228181
2022-03-20   -1.676719
2022-03-21    0.244249
2022-03-22    1.776478
2022-03-23   -0.185076
2022-03-24   -0.216682
2022-03-25    0.624784
Freq: D, Name: A, dtype: float64
=====
```

```
df["A"]
2022-03-19    0.228181
2022-03-20   -1.676719
2022-03-21    0.244249
2022-03-22    1.776478
2022-03-23   -0.185076
2022-03-24   -0.216682
2022-03-25    0.624784
Freq: D, Name: A, dtype: float64
```

Selecting via `[]`, which slicing the rows.

```
[ ]: df[0:3]
```

```
[ ]:
      A          B          C          D
2022-03-19  0.228181  0.411725 -0.888684 -0.907902
2022-03-20 -1.676719  1.090258 -1.966077  0.909066
2022-03-21  0.244249 -0.256577 -0.099083 -0.169163
```

```
[ ]: df["20220319":"20220321"]
```

```
[ ]:
      A          B          C          D
2022-03-19  0.228181  0.411725 -0.888684 -0.907902
2022-03-20 -1.676719  1.090258 -1.966077  0.909066
2022-03-21  0.244249 -0.256577 -0.099083 -0.169163
```

1.4 Selection by label

For getting a cross section using a label.

```
[ ]: df.loc[dates[0]]
```

```
[ ]: A    0.228181  
     B    0.411725  
     C   -0.888684  
     D   -0.907902  
     Name: 2022-03-19 00:00:00, dtype: float64
```

Selecting on a multi-axis by label.

```
[ ]: df.loc[:, ["A", "B"]]
```

```
[ ]:           A      B  
2022-03-19  0.228181  0.411725  
2022-03-20 -1.676719  1.090258  
2022-03-21  0.244249 -0.256577  
2022-03-22  1.776478  0.235238  
2022-03-23 -0.185076 -0.964747  
2022-03-24 -0.216682  0.182571  
2022-03-25  0.624784 -0.974044
```

Showing label slicing, both endpoints are included.

```
[ ]: df.loc["20220319":"20220321", ["C", "D"]]
```

```
[ ]:           C      D  
2022-03-19 -0.888684 -0.907902  
2022-03-20 -1.966077  0.909066  
2022-03-21 -0.099083 -0.169163
```

Reduction in the dimensions of returned object.

```
[ ]: df.loc["20220324", ["B", "C"]]
```

```
[ ]: B    0.182571  
     C    0.885816  
     Name: 2022-03-24 00:00:00, dtype: float64
```

For getting a scalar value.

```
[ ]: df.loc[dates[1], "A"]
```

```
[ ]: -1.6767189854604438
```

For getting fast access to a scalar (equivalent to the prior method).

```
[ ]: df.at[dates[1], "A"]
```

```
[ ]: -1.6767189854604438
```

1.5 Selection by position

Select via the position of the passed integers.

```
[ ]: df.iloc[3]  
# equivalent to df.loc[dates[3]]
```

```
[ ]: A    1.776478  
     B    0.235238  
     C   -0.093981  
     D    0.867664  
     Name: 2022-03-22 00:00:00, dtype: float64
```

By integer slices, acting similar to Numpy/Python.

```
[ ]: df.iloc[3:5, 0:2]  
# equivalent to df.loc["20220322":"20220323", ["A", "B"]]
```

```
[ ]:           A      B  
2022-03-22  1.776478  0.235238  
2022-03-23 -0.185076 -0.964747
```

By list of integer position locations, similar to the Numpy/Python style.

```
[ ]: df.iloc[[1, 2, 4], [0, 2]]  
# equivalent to df.loc["20220320":"20220323", ["B", "C"]]
```

```
[ ]:           A      C  
2022-03-20 -1.676719 -1.966077  
2022-03-21  0.244249 -0.099083  
2022-03-23 -0.185076 -0.542850
```

For slicing rows explicitly.

```
[ ]: df.iloc[1:3, :]
```

```
[ ]:           A      B      C      D  
2022-03-20 -1.676719  1.090258 -1.966077  0.909066  
2022-03-21  0.244249 -0.256577 -0.099083 -0.169163
```

For slicing columns explicitly.

```
[ ]: df.iloc[:, 1:3]
```



```
[ ]:
      B      C
2022-03-19  0.411725 -0.888684
2022-03-20  1.090258 -1.966077
2022-03-21 -0.256577 -0.099083
2022-03-22  0.235238 -0.093981
2022-03-23 -0.964747 -0.542850
2022-03-24  0.182571  0.885816
2022-03-25 -0.974044  1.100455
```

For getting a value explicitly.

```
[ ]: df.iloc[1, 1]
```

```
[ ]: 1.0902580667797517
```

For getting fast access to a scalar (equivalent to the prior method).

```
[ ]: df.iat[1, 1]
```

```
[ ]: 1.0902580667797517
```

1.6 Boolean indexing

Using a single column's values to select data.

```
[ ]: df[df["A"] > 0]
```

```
[ ]:
      A      B      C      D
2022-03-19  0.228181  0.411725 -0.888684 -0.907902
2022-03-21  0.244249 -0.256577 -0.099083 -0.169163
2022-03-22  1.776478  0.235238 -0.093981  0.867664
2022-03-25  0.624784 -0.974044  1.100455  0.893449
```

Selecting values from a **DataFrame** where a boolean condition is met.

```
[ ]: df[df > 0]
```

```
[ ]:
      A      B      C      D
2022-03-19  0.228181  0.411725    NaN    NaN
2022-03-20    NaN  1.090258    NaN  0.909066
2022-03-21  0.244249    NaN    NaN    NaN
2022-03-22  1.776478  0.235238    NaN  0.867664
2022-03-23    NaN    NaN    NaN    NaN
2022-03-24    NaN  0.182571  0.885816    NaN
2022-03-25  0.624784    NaN  1.100455  0.893449
```

Using `isin()` method for filtering.

```
[ ]: df_copy = df.copy()
df_copy["E"] = ["one", "two", "three", "four", "five", "six", "seven"]
df_copy
```

```
[ ]:
      A          B          C          D          E
2022-03-19  0.228181  0.411725 -0.888684 -0.907902    one
2022-03-20 -1.676719  1.090258 -1.966077  0.909066    two
2022-03-21  0.244249 -0.256577 -0.099083 -0.169163   three
2022-03-22  1.776478  0.235238 -0.093981  0.867664    four
2022-03-23 -0.185076 -0.964747 -0.542850 -0.283645    five
2022-03-24 -0.216682  0.182571  0.885816 -0.151794    six
2022-03-25  0.624784 -0.974044  1.100455  0.893449   seven
```

```
[ ]: df_copy[df_copy["E"].isin(["three", "six"])]
```

```
[ ]:
      A          B          C          D          E
2022-03-21  0.244249 -0.256577 -0.099083 -0.169163   three
2022-03-24 -0.216682  0.182571  0.885816 -0.151794    six
```

1.7 Setting

Setting a new column automatically aligns the data by the indexes.

```
[ ]: s1 = pd.Series(list(range(1, 8)), index=pd.date_range("20220319", periods=7))
s1
```

```
[ ]: 2022-03-19    1
      2022-03-20    2
      2022-03-21    3
      2022-03-22    4
      2022-03-23    5
      2022-03-24    6
      2022-03-25    7
Freq: D, dtype: int64
```

```
[ ]: df["F"] = s1
df
```

```
[ ]:
      A          B          C          D          F
2022-03-19  0.228181  0.411725 -0.888684 -0.907902    1
2022-03-20 -1.676719  1.090258 -1.966077  0.909066    2
2022-03-21  0.244249 -0.256577 -0.099083 -0.169163    3
2022-03-22  1.776478  0.235238 -0.093981  0.867664    4
2022-03-23 -0.185076 -0.964747 -0.542850 -0.283645    5
2022-03-24 -0.216682  0.182571  0.885816 -0.151794    6
2022-03-25  0.624784 -0.974044  1.100455  0.893449    7
```

Setting values by label.

```
[ ]: df.at[dates[0], "A"] = 0
df
```

```
[ ]:
```

	A	B	C	D	F
2022-03-19	0.000000	0.411725	-0.888684	-0.907902	1
2022-03-20	-1.676719	1.090258	-1.966077	0.909066	2
2022-03-21	0.244249	-0.256577	-0.099083	-0.169163	3
2022-03-22	1.776478	0.235238	-0.093981	0.867664	4
2022-03-23	-0.185076	-0.964747	-0.542850	-0.283645	5
2022-03-24	-0.216682	0.182571	0.885816	-0.151794	6
2022-03-25	0.624784	-0.974044	1.100455	0.893449	7

Setting values by position.

```
[ ]: df.iat[0, 1] = 1
df
```

```
[ ]:
```

	A	B	C	D	F
2022-03-19	0.000000	1.000000	-0.888684	-0.907902	1
2022-03-20	-1.676719	1.090258	-1.966077	0.909066	2
2022-03-21	0.244249	-0.256577	-0.099083	-0.169163	3
2022-03-22	1.776478	0.235238	-0.093981	0.867664	4
2022-03-23	-0.185076	-0.964747	-0.542850	-0.283645	5
2022-03-24	-0.216682	0.182571	0.885816	-0.151794	6
2022-03-25	0.624784	-0.974044	1.100455	0.893449	7

Setting by assigning with a Numpy array

```
[ ]: df.loc[:, "D"] = np.array(5 * len(df))
df
```

```
[ ]:
```

	A	B	C	D	F
2022-03-19	0.000000	1.000000	-0.888684	35	1
2022-03-20	-1.676719	1.090258	-1.966077	35	2
2022-03-21	0.244249	-0.256577	-0.099083	35	3
2022-03-22	1.776478	0.235238	-0.093981	35	4
2022-03-23	-0.185076	-0.964747	-0.542850	35	5
2022-03-24	-0.216682	0.182571	0.885816	35	6
2022-03-25	0.624784	-0.974044	1.100455	35	7

1.8 Missing data

Pandas primarily uses the value `np.nan` to represent missing data. It is by default not included in computation.

Reindexing allows you to change/add/delete the index on a sepecified axis. This returns a copy of data.

```
[ ]: df1 = df.reindex(index=dates[0:4], columns=list(df.columns) + ["E"])
df1.loc[dates[0]:dates[1], "E"] = 1
df1
```

```
[ ]:
      A          B          C  D  F  E
2022-03-19  0.000000  1.000000 -0.888684  35  1  1.0
2022-03-20 -1.676719  1.090258 -1.966077  35  2  1.0
2022-03-21  0.244249 -0.256577 -0.099083  35  3  NaN
2022-03-22  1.776478  0.235238 -0.093981  35  4  NaN
```

To drop any rows that have missing data.

```
[ ]: df1.dropna(how="any")
```

```
[ ]:
      A          B          C  D  F  E
2022-03-19  0.000000  1.000000 -0.888684  35  1  1.0
2022-03-20 -1.676719  1.090258 -1.966077  35  2  1.0
```

Filling missing data.

```
[ ]: df1.fillna(value=3)
```

```
[ ]:
      A          B          C  D  F  E
2022-03-19  0.000000  1.000000 -0.888684  35  1  1.0
2022-03-20 -1.676719  1.090258 -1.966077  35  2  1.0
2022-03-21  0.244249 -0.256577 -0.099083  35  3  3.0
2022-03-22  1.776478  0.235238 -0.093981  35  4  3.0
```

To get the boolean mask where values are nan.

```
[ ]: pd.isna(df1)
```

```
[ ]:
      A          B          C  D  F  E
2022-03-19  False  False  False  False  False  False
2022-03-20  False  False  False  False  False  False
2022-03-21  False  False  False  False  False  True
2022-03-22  False  False  False  False  False  True
```

1.9 Operations

1.9.1 Stats

Operations in general *exclude* missing data.

Perform a descriptive statistic.

```
[ ]: df.mean()
```

```
[ ]: A      0.081005  
     B      0.044671  
     C     -0.229200  
     D     35.000000  
     F      4.000000  
     dtype: float64
```

Same operation on the other axis.

```
[ ]: df.mean(1)
```

```
[ ]: 2022-03-19      7.222263  
     2022-03-20      6.889492  
     2022-03-21      7.577718  
     2022-03-22      8.183547  
     2022-03-23      7.661465  
     2022-03-24      8.370341  
     2022-03-25      8.550239  
     Freq: D, dtype: float64
```

Operation with objects that have different dimensionality and need alignment. In addition, `pandas` automatically broadcast along the specified dimension.

```
[ ]: s = pd.Series([1, 3, 5, np.nan, 6, 8, 11], index=dates).shift(2)  
     s
```

```
[ ]: 2022-03-19      NaN  
     2022-03-20      NaN  
     2022-03-21      1.0  
     2022-03-22      3.0  
     2022-03-23      5.0  
     2022-03-24      NaN  
     2022-03-25      6.0  
     Freq: D, dtype: float64
```

```
[ ]: df.sub(s, axis="index")
```

```
[ ]:           A      B      C      D      F  
     2022-03-19      NaN      NaN      NaN      NaN      NaN  
     2022-03-20      NaN      NaN      NaN      NaN      NaN  
     2022-03-21 -0.755751 -1.256577 -1.099083 34.0 2.0  
     2022-03-22 -1.223522 -2.764762 -3.093981 32.0 1.0  
     2022-03-23 -5.185076 -5.964747 -5.542850 30.0 0.0  
     2022-03-24      NaN      NaN      NaN      NaN      NaN  
     2022-03-25 -5.375216 -6.974044 -4.899545 29.0 1.0
```

1.9.2 Apply

Applying functions to the data.

```
[ ]: df
```

```
[ ]:
      A      B      C  D  F
2022-03-19  0.000000  1.000000 -0.888684  35  1
2022-03-20 -1.676719  1.090258 -1.966077  35  2
2022-03-21  0.244249 -0.256577 -0.099083  35  3
2022-03-22  1.776478  0.235238 -0.093981  35  4
2022-03-23 -0.185076 -0.964747 -0.542850  35  5
2022-03-24 -0.216682  0.182571  0.885816  35  6
2022-03-25  0.624784 -0.974044  1.100455  35  7
```

```
[ ]: df.apply(np.cumsum)
# equivalent to df.apply(np.cumsum, axis=0)
```

```
[ ]:
      A      B      C  D  F
2022-03-19  0.000000  1.000000 -0.888684  35  1
2022-03-20 -1.676719  2.090258 -2.854761  70  3
2022-03-21 -1.432470  1.833681 -2.953844  105  6
2022-03-22  0.344008  2.068919 -3.047825  140  10
2022-03-23  0.158932  1.104172 -3.590675  175  15
2022-03-24 -0.057751  1.286744 -2.704859  210  21
2022-03-25  0.567034  0.312700 -1.604403  245  28
```

```
[ ]: df.apply(np.cumsum, axis=1)
```

```
[ ]:
      A      B      C      D      F
2022-03-19  0.000000  1.000000  0.111316  35.111316  36.111316
2022-03-20 -1.676719 -0.586461 -2.552538  32.447462  34.447462
2022-03-21  0.244249 -0.012328 -0.111411  34.888589  37.888589
2022-03-22  1.776478  2.011716  1.917735  36.917735  40.917735
2022-03-23 -0.185076 -1.149823 -1.692673  33.307327  38.307327
2022-03-24 -0.216682 -0.034111  0.851705  35.851705  41.851705
2022-03-25  0.624784 -0.349260  0.751195  35.751195  42.751195
```

```
[ ]: df.apply(lambda x: x.max() - x.min())
```

```
[ ]: A    3.453197
      B    2.064302
      C    3.066532
      D    0.000000
      F    6.000000
      dtype: float64
```

```
[ ]: df.apply(lambda x: x.max() - x.min(), axis=1)
```

```
[ ]: 2022-03-19    35.888684
      2022-03-20    36.966077
      2022-03-21    35.256577
      2022-03-22    35.093981
      2022-03-23    35.964747
      2022-03-24    35.216682
      2022-03-25    35.974044
      Freq: D, dtype: float64
```

1.9.3 Histogramming

```
[ ]: s = pd.Series(np.random.randint(0, 7, size=10))
      s
```

```
[ ]: 0    6
      1    0
      2    4
      3    4
      4    1
      5    2
      6    0
      7    6
      8    1
      9    2
      dtype: int32
```

```
[ ]: s.value_counts()
```

```
[ ]: 6    2
      0    2
      4    2
      1    2
      2    2
      dtype: int64
```

1.9.4 String methods

Series is equipped with a set of string processing methods in the `str` attribute that make it easy to operate on each element of the array, as in the code snippet below. Note that pattern-matching in `str` generally uses **regular expressions** by default (and in some cases always uses them).

```
[ ]: s = pd.Series(["A", "B", "C", "D", "Aaba", "Baca", np.nan, "CABA", "dog", "cat"])
      s
```

```
[ ]: 0      A
      1      B
      2      C
      3      D
      4    Aaba
      5    Baca
      6     NaN
      7    CABA
      8    dog
      9    cat
      dtype: object
```

```
[ ]: s.str.lower()
```

```
[ ]: 0      a
      1      b
      2      c
      3      d
      4    aaba
      5    baca
      6     NaN
      7    caba
      8    dog
      9    cat
      dtype: object
```

1.9.5 Merge

Concat Pandas provides various facilities for easily combining together **Series** and **DataFrame** objects with various kinds of set logic for the indexes and relational algebra functionality in the case of join/merge-type operations.

Concatenating pandas objects together with `concat()`.

```
[ ]: df = pd.DataFrame(np.random.randn(10, 4))
      df
```

```
[ ]:      0      1      2      3
0 -1.724109  0.411101 -0.328557  0.250358
1 -1.254807  1.463416  1.133394 -1.834136
2 -0.607070 -1.460862 -0.433120 -0.098387
3 -0.236818  0.420965 -0.356506  0.323899
4 -0.603135  1.060861  0.405774 -0.511619
5 -0.582767 -0.639265 -0.033607 -0.469928
6 -0.339338  0.181223  0.228185  0.822008
7  0.213306 -0.992878  0.517584  0.304642
8 -1.147936 -0.083162 -0.256240  0.034513
```



```
9 0.266978 -0.413701 0.114007 0.177897
```

```
[ ]: # break it into pieces
pieces = [df[:3], df[3:7], df[7:]]
pieces
```

```
[ ]: [
      0      1      2      3
0 -1.724109  0.411101 -0.328557  0.250358
1 -1.254807  1.463416  1.133394 -1.834136
2 -0.607070 -1.460862 -0.433120 -0.098387,
      0      1      2      3
3 -0.236818  0.420965 -0.356506  0.323899
4 -0.603135  1.060861  0.405774 -0.511619
5 -0.582767 -0.639265 -0.033607 -0.469928
6 -0.339338  0.181223  0.228185  0.822008,
      0      1      2      3
7  0.213306 -0.992878  0.517584  0.304642
8 -1.147936 -0.083162 -0.256240  0.034513
9  0.266978 -0.413701  0.114007  0.177897]
```

```
[ ]: pd.concat(pieces)
```

```
[ ]:
      0      1      2      3
0 -1.724109  0.411101 -0.328557  0.250358
1 -1.254807  1.463416  1.133394 -1.834136
2 -0.607070 -1.460862 -0.433120 -0.098387
3 -0.236818  0.420965 -0.356506  0.323899
4 -0.603135  1.060861  0.405774 -0.511619
5 -0.582767 -0.639265 -0.033607 -0.469928
6 -0.339338  0.181223  0.228185  0.822008
7  0.213306 -0.992878  0.517584  0.304642
8 -1.147936 -0.083162 -0.256240  0.034513
9  0.266978 -0.413701  0.114007  0.177897
```

Adding a column to a `DataFrame` is relatively fast. However, adding a row requires a copy, and may be expensive. I recommend passing a pre-built list of records to the `DataFrame` constructor instead of building a `DataFrame` by iteratively appending records to it.

1.9.6 Join

SQL style merges.

```
[ ]: left = pd.DataFrame(
    {
        "keys": ["foo", "bar"],
        "lval": [1, 2]
```

```

    }
)
right = pd.DataFrame(
    {
        "keys": ["foo", "bar"],
        "rval": [4, 5]
    }
)

```

```
[ ]: left
```

```
[ ]:  keys  lval
0  foo     1
1  bar     2
```

```
[ ]: right
```

```
[ ]:  keys  rval
0  foo     4
1  bar     5
```

```
[ ]: pd.merge(left, right, on="keys")
```

```
[ ]:  keys  lval  rval
0  foo     1     4
1  bar     2     5
```

1.10 Grouping

By "grouping" we are referring to a process involving one or more of the following steps:

- **Splitting** the data into groups based on some criteria.
- **Applying** a function to each group independently
- **Combining** the results into a data structure

```
[ ]: df = pd.DataFrame(
    {
        "A": ["foo", "bar", "foo", "bar", "foo", "bar", "foo", "foo"],
        "B": ["one", "one", "two", "three", "two", "two", "one", "three"],
        "C": np.random.randn(8),
        "D": np.random.randn(8),
    }
)
df

```

```
[ ]:   A      B      C      D
0  foo  one  0.338328  0.407254
```

```

1 bar    one  1.277655  0.088752
2 foo    two -0.725769  1.240329
3 bar   three -0.326743 -0.559055
4 foo    two -0.089508 -0.495114
5 bar    two -0.082540 -0.471225
6 foo    one -1.126816 -0.151158
7 foo   three  2.010950  1.234230

```

Grouping and then applying the `sum()` function to the resulting group.

```
[ ]: df.groupby("A").sum()
```

```
[ ]:
      C      D
A
bar  0.868371 -0.941528
foo  0.407185  2.235541

```

Grouping by multiple columns forms a hierarchical index, and again we can apply the `sum()` function.

```
[ ]: df.groupby(["A", "B"]).sum()
```

```
[ ]:
      C      D
A  B
bar one  1.277655  0.088752
    three -0.326743 -0.559055
    two  -0.082540 -0.471225
foo one  -0.788489  0.256096
    three  2.010950  1.234230
    two  -0.815277  0.745215

```

1.11 Reshaping

1.11.1 Stack

```
[ ]: tuples = list(
    zip(
        *[
            ["bar", "bar", "baz", "baz", "foo", "foo", "qux", "qux"],
            ["one", "two", "one", "two", "one", "two", "one", "two"],
        ]
    )
)
index = pd.MultiIndex.from_tuples(tuples, names=["first", "second"])
df = pd.DataFrame(np.random.randn(8, 2), index=index, columns=["A", "B"])
df

```

```
[ ]:
      first second      A      B
bar   one      0.530897 -0.825297
      two     -0.296452 -0.448199
baz   one      0.953647 -0.304044
      two     -0.871343 -1.245235
foo   one     -0.405061 -0.422861
      two     -0.973405 -0.167809
qux   one      0.399787 -0.243971
      two      2.014824 -1.222681
```

The `stack()` method "compress" a level in the **DataFrame's** columns.

```
[ ]: df2 = df[:4]
      df2
```

```
[ ]:
      first second      A      B
bar   one      0.530897 -0.825297
      two     -0.296452 -0.448199
baz   one      0.953647 -0.304044
      two     -0.871343 -1.245235
```

```
[ ]: stacked = df2.stack()
      stacked
```

```
[ ]: first  second
bar    one      A    0.530897
      two      A   -0.296452
      two      B   -0.448199
baz    one      A    0.953647
      two      B   -0.304044
      two      A   -0.871343
      two      B   -1.245235
dtype: float64
```

With a "stacked" DataFrame or Series (having a **MultiIndex** as the **Index**), the inverse operation of `stack()` is `unstack()`, which by default unstacks the **last** level.

```
[ ]: stacked.unstack()
```

```
[ ]:
      first second      A      B
bar   one      0.530897 -0.825297
      two     -0.296452 -0.448199
baz   one      0.953647 -0.304044
      two     -0.871343 -1.245235
```

```
[ ]: stacked.unstack(1)
```

```
[ ]: second      one      two
first
bar   A  0.530897 -0.296452
      B -0.825297 -0.448199
baz   A  0.953647 -0.871343
      B -0.304044 -1.245235
```

```
[ ]: stacked.unstack(0)
```

```
[ ]: first      bar      baz
second
one   A  0.530897  0.953647
      B -0.825297 -0.304044
two   A -0.296452 -0.871343
      B -0.448199 -1.245235
```

1.12 Pivot tables

```
[ ]: df = pd.DataFrame(
    {
        "A": ["one", "one", "two", "three"] * 3,
        "B": ["A", "B", "C"] * 4,
        "C": ["foo", "foo", "foo", "bar", "bar", "bar"] * 2,
        "D": np.random.randn(12),
        "E": np.random.randn(12),
    }
)
df
```

```
[ ]:
   A  B  C      D      E
0  one A  foo  1.488596 -0.117750
1  one B  foo  0.363735 -0.567175
2  two C  foo -1.184895 -0.880365
3 three A  bar  1.710122 -0.495467
4  one B  bar  1.319943 -0.564924
5  one C  bar  0.061341 -0.416826
6  two A  foo -0.890777  1.249897
7 three B  foo -0.062211 -1.280127
8  one C  foo -0.786301  1.806379
9  one A  bar  0.341206 -1.359112
10 two B  bar  1.131091  0.694300
11 three C  bar -0.550187 -0.010789
```

We can produce pivot tables from this data very easily.

```
[ ]: pd.pivot_table(df, values="D", index=["A", "B"], columns="C")
```

```
[ ]: C          bar          foo
A      B
one  A  0.341206  1.488596
     B  1.319943  0.363735
     C  0.061341 -0.786301
three A  1.710122         NaN
     B         NaN -0.062211
     C -0.550187         NaN
two  A         NaN -0.890777
     B  1.131091         NaN
     C         NaN -1.184895
```

1.13 Time series

Pandas has simple, powerful, and efficient functionality for performing resampling operations during frequency conversion (e.g., converting secondly data into 5-minutely data). This is extremely common in, but not limited to, financial applications.

```
[ ]: rng = pd.date_range("1/1/2022", periods=100, freq="S")
     ts = pd.Series(np.random.randint(0, 500, len(rng)), index=rng)
     ts
```

```
[ ]: 2022-01-01 00:00:00    251
     2022-01-01 00:00:01    489
     2022-01-01 00:00:02    348
     2022-01-01 00:00:03    206
     2022-01-01 00:00:04     42
     ...
     2022-01-01 00:01:35    408
     2022-01-01 00:01:36    126
     2022-01-01 00:01:37    483
     2022-01-01 00:01:38     24
     2022-01-01 00:01:39    144
     Freq: S, Length: 100, dtype: int32
```

```
[ ]: ts.resample("5Min").sum()
```

```
[ ]: 2022-01-01    25972
     Freq: 5T, dtype: int32
```

Time zone representation

```
[ ]: rng = pd.date_range("1/1/2022 00:00", periods=5, freq="D")
     ts = pd.Series(np.random.randn(len(rng)), rng)
     ts
```

```
[ ]: 2022-01-01    -0.727610
     2022-01-02     0.653522
     2022-01-03    -0.104993
     2022-01-04    -0.798697
     2022-01-05     1.556841
     Freq: D, dtype: float64
```

```
[ ]: ts_utc = ts.tz_localize("UTC")
     ts_utc
```

```
[ ]: 2022-01-01 00:00:00+00:00    -0.727610
     2022-01-02 00:00:00+00:00     0.653522
     2022-01-03 00:00:00+00:00    -0.104993
     2022-01-04 00:00:00+00:00    -0.798697
     2022-01-05 00:00:00+00:00     1.556841
     Freq: D, dtype: float64
```

Converting to another time zone.

```
[ ]: ts_utc.tz_convert("US/Eastern")
```

```
[ ]: 2021-12-31 19:00:00-05:00    -0.727610
     2022-01-01 19:00:00-05:00     0.653522
     2022-01-02 19:00:00-05:00    -0.104993
     2022-01-03 19:00:00-05:00    -0.798697
     2022-01-04 19:00:00-05:00     1.556841
     Freq: D, dtype: float64
```

Converting between time span representations.

```
[ ]: rng = pd.date_range("1/1/2022", periods=5, freq="M")
     ts = pd.Series(np.random.randn(len(rng)), index=rng)
     ts
```

```
[ ]: 2022-01-31    -0.524220
     2022-02-28    -0.768486
     2022-03-31    -1.456069
     2022-04-30    -1.391087
     2022-05-31    -0.413119
     Freq: M, dtype: float64
```

```
[ ]: ps = ts.to_period()
     ps
```

```
[ ]: 2022-01    -0.524220
      2022-02    -0.768486
      2022-03    -1.456069
      2022-04    -1.391087
      2022-05    -0.413119
      Freq: M, dtype: float64
```

```
[ ]: ps.to_timestamp()
```

```
[ ]: 2022-01-01    -0.524220
      2022-02-01    -0.768486
      2022-03-01    -1.456069
      2022-04-01    -1.391087
      2022-05-01    -0.413119
      Freq: MS, dtype: float64
```

Converting between period and timestamp enables some convenient arithmetic functions to be used. In the following example, I convert a quarterly frequency with year ending in November to 9am of the end of the month following the quarter end.

```
[ ]: prng = pd.period_range("1990Q1", "2000Q4", freq="Q-NOV")
      ts = pd.Series(np.random.randn(len(prng)), prng)
      ts.index = (prng.asfreq("M", "e") + 1).asfreq("H", "s") + 9
      ts.head()
```

```
[ ]: 1990-03-01 09:00    0.915812
      1990-06-01 09:00    0.924666
      1990-09-01 09:00    0.366394
      1990-12-01 09:00    0.370943
      1991-03-01 09:00    0.172963
      Freq: H, dtype: float64
```

1.14 Categoricals

Pandas can include categorical data in a **DataFrame**.

```
[ ]: df = pd.DataFrame(
      {
          "id": list(range(1, 7, 1)),
          "raw_grade": ["a", "a", "b", "a", "a", "e"],
      }
    )
```

Converting the grades to a categorical data type.

```
[ ]: df["grade"] = df["raw_grade"].astype("category")
      df["grade"]
```



```
[ ]: 0    a
      1    a
      2    b
      3    a
      4    a
      5    e
      Name: grade, dtype: category
      Categories (3, object): ['a', 'b', 'e']
```

Rename the categories to more meaningful names (assigning to `Series.cat.categories()` is in place!).

```
[ ]: df["grade"].cat.categories = ["very good", "good", "very bad"]
      df["grade"]
```

```
[ ]: 0    very good
      1    very good
      2         good
      3    very good
      4    very good
      5    very bad
      Name: grade, dtype: category
      Categories (3, object): ['very good', 'good', 'very bad']
```

Reorder the categories and simultaneously add the missing categories (methods under `Series.cat()` return a new `Series` by default).

```
[ ]: df["grade"] = df["grade"].cat.set_categories(
      ["very bad", "bad", "medium", "good", "very good"]
      )
      df["grade"]
```

```
[ ]: 0    very good
      1    very good
      2         good
      3    very good
      4    very good
      5    very bad
      Name: grade, dtype: category
      Categories (5, object): ['very bad', 'bad', 'medium', 'good', 'very good']
```

Sorting in per order in the categories, not lexical order.

```
[ ]: df.sort_values(by="grade", ascending=False)
```

```
[ ]:   id raw_grade  grade
      0    1      a  very good
      1    2      a  very good
```

3	4	a	very good
4	5	a	very good
2	3	b	good
5	6	e	very bad

Grouping by a categorical column also shows empty categories.

```
[ ]: df.groupby("grade").size()
```

```
[ ]: grade
very bad    1
bad         0
medium      0
good        1
very good   4
dtype: int64
```

1.15 Plotting


We use the standard convention for referencing the matplotlib API.

```
[ ]: import matplotlib.pyplot as plt
plt.close("all")
```

The `close()` method is used to close a figure window.


```
[ ]: ts = pd.Series(np.random.randn(1000), index=pd.date_range("1/1/2022",
↪ periods=1000))
ts.plot()
```

```
[ ]: <AxesSubplot:>
```



```
[ ]: ts_cumsum = ts.cumsum()  
ts_cumsum.plot()
```

```
[ ]: <AxesSubplot:>
```



If running under Jupyter Notebook, the plot will appear on `plot()`. Otherwise use `matplotlib.pyplot.show` to show it or `matplotlib.pyplot.savefig` to write it into a file.


```
[ ]: plt.show()
```

On a `DataFrame`, the `plot()` method is a convenience to plot all of the columns with labels.

```
[ ]: df = pd.DataFrame(  
    np.random.randn(1000, 4),  
    index=ts.index,  
    columns=list("ABCD")  
)  
df = df.cumsum()  
plt.figure()  
df.plot()  
plt.legend(loc="best")
```

```
[ ]: <matplotlib.legend.Legend at 0x1e0fd34ec10>
```

<Figure size 432x288 with 0 Axes>



1.16 Getting data in/out

1.16.1 CSV

Writing to a .csv file

```
[ ]: df.to_csv("files/foo.csv")
```

Reading from a .csv file

```
[ ]: pd.read_csv("files/foo.csv")
```

```
[ ]:      Unnamed: 0      A      B      C      D
0    2022-01-01  0.091463  1.673725 -0.746928  0.857382
1    2022-01-02  0.767879  2.172299 -1.809359  0.242449
2    2022-01-03 -0.188135  3.627265 -0.954293  0.558939
3    2022-01-04 -0.568480  3.282168 -2.156548 -0.169787
4    2022-01-05 -0.031601  1.061619 -3.950347  0.062531
..      ...      ...      ...      ...      ...
995  2024-09-22 -15.264482 -29.471355 -21.530166  24.533283
996  2024-09-23 -14.416862 -29.069515 -22.524112  23.382795
997  2024-09-24 -14.521215 -28.743473 -21.557216  21.931433
998  2024-09-25 -15.213408 -28.746767 -19.732183  23.176940
999  2024-09-26 -16.628631 -28.289745 -20.632969  24.248229
```

[1000 rows x 5 columns]

1.16.2 HDF5

Reading and writing to HDF5Stores.

Writing to a HDF5 Store.

```
[ ]: # !pip install tables
df.to_hdf("files/foo.h5", "df")
```

Reading from HDF5 Store.

```
[ ]: pd.read_hdf("files/foo.h5", "df")
```

```
[ ]:      A      B      C      D
2022-01-01  0.091463  1.673725 -0.746928  0.857382
2022-01-02  0.767879  2.172299 -1.809359  0.242449
2022-01-03 -0.188135  3.627265 -0.954293  0.558939
2022-01-04 -0.568480  3.282168 -2.156548 -0.169787
2022-01-05 -0.031601  1.061619 -3.950347  0.062531
...      ...      ...      ...      ...
```

```

2024-09-22 -15.264482 -29.471355 -21.530166 24.533283
2024-09-23 -14.416862 -29.069515 -22.524112 23.382795
2024-09-24 -14.521215 -28.743473 -21.557216 21.931433
2024-09-25 -15.213408 -28.746767 -19.732183 23.176940
2024-09-26 -16.628631 -28.289745 -20.632969 24.248229

```

[1000 rows x 4 columns]

1.16.3 Excel

Reading and writing to MS Excel.

Writing to an excel file

```
[ ]: df.to_excel("files/foo.xlsx", sheet_name="Sheet1")
```

Reading from an excel file.

```
[ ]: pd.read_excel("files/foo.xlsx", "Sheet1", index_col=None, na_values=["NA"])
```

```
[ ]:
   Unnamed: 0      A      B      C      D
0  2022-01-01  0.091463  1.673725 -0.746928  0.857382
1  2022-01-02  0.767879  2.172299 -1.809359  0.242449
2  2022-01-03 -0.188135  3.627265 -0.954293  0.558939
3  2022-01-04 -0.568480  3.282168 -2.156548 -0.169787
4  2022-01-05 -0.031601  1.061619 -3.950347  0.062531
..      ...      ...      ...      ...
995 2024-09-22 -15.264482 -29.471355 -21.530166 24.533283
996 2024-09-23 -14.416862 -29.069515 -22.524112 23.382795
997 2024-09-24 -14.521215 -28.743473 -21.557216 21.931433
998 2024-09-25 -15.213408 -28.746767 -19.732183 23.176940
999 2024-09-26 -16.628631 -28.289745 -20.632969 24.248229

```

[1000 rows x 5 columns]